Evaluating Tensor Cores

John Carlsson, Cyprien Toulon de Coutex
Department of Computer Science
University of Salerno
Italy

Abstract— This project was aimed at evaluating Tensor cores in cuda through experiemnts to get a better understanding off high performance computing through different methods. With the ultimate goal to compare the time to calculate a matrix multiplication addition. At the start of the project, we were fairly certain that the fastest algorithm would be the Tensor cores, followed by Cuda with shared memory,Cuda global memory, and at the bottom, an optimized CPU version. However, this was not the case. The results show that Cuda shared memory is the fastest with global memory just behind. For small matrices tensors appear superior.

I. INTRODUCTION

Over the past decade, the increasing demand for deep learning and AI applications has fueled the need for highly efficient tensor operations. Traditional GPU architectures are powerful but are often limited in their ability to fully exploit the potential of tensor computations due to the reliance on higher-precision floating-point arithmetic. Tensor Cores address this limitation by leveraging mixed-precision arithmetic, combining high throughput with reduced precision.

Tensor Cores, introduced in NVIDIA's Volta architecture and further enhanced in subsequent architectures such as Turing and Ampere, have revolutionized the performance of tensor computations on GPUs. These specialized hardware units offer significant speedups by providing native support for low-precision tensor operations, specifically matrix multiplications and convolutions. The purpose of this project report is to present an evaluation of Tensor Cores, exploring their capabilities compared to CPU and GPU based computation.

In this project, we aim to evaluate Tensor Cores using the Matrix Multiplication Accumulate (MMA) with various workloads and compare the results to the CPU and GPU performances.

By conducting this comprehensive evaluation, we aim to provide insights into the capabilities and limitations of Tensor Cores. Additionally, the findings from this study will contribute to the broader understanding of GPU acceleration techniques for tensor operations and their potential impact on high-performance computing.

A. Motivation

The MMA operation plays a critical role in accelerating various computational tasks. It involves three matrices and is fundamental within many computational domains, such as deep learning, scientific simulations, and image processing. However, performing MMA efficiently and at scale is difficult due to the intensity of the operation.

The difficulties that arise when trying to implement a fast version of MMA are due to a combination of factors including memory access patterns, data dependencies, computational intensity, matrix sizes, hardware limitations, and optimization trade-offs.

The objective of this paper is to provide valuable insights for optimizing MMA computations and facilitate the adoption of Tensor Cores as a powerful tool for accelerating MMA operations in comparison to CPU and GPU. Our contributions will be as follows:

- We will measure the speedup achieved by Tensor Cores measured in actual seconds and explore how performance scales. with the size of the matrices.
- We will compare the performance of Tensor Cores with CPU and GPU implementations of MMA using different matrix sizes and precision settings.
- We will analyze the impact of data precision on the performance of Tensor Cores and investigate the trade-offs between accuracy and computation time.
- We will explore optimization techniques for maximizing the utilization of Tensor Cores and achieving optimal performance.
- We will provide insights into the limitations and challenges of Tensor Cores, including memory constraints, data dependencies, and scalability issues.

II. BACKGROUND

A. Tensor Cores

Fig. 1. SM architecture.[?]

Tensor Cores are specialized hardware units introduced in NVIDIA GPUs that are designed to accelerate matrix

computations, specifically matrix multiplications and convolutions. They leverage mixed-precision arithmetic, combining high throughput with reduced precision to achieve significant speedups in tensor operations.

The key feature of Tensor Cores is their ability to perform mixed-precision operations using half-precision floating-point (FP16) data types. By using FP16, Tensor Cores can process a larger number of elements simultaneously, leading to improved throughput. Additionally, Tensor Cores support automatic data type conversion, allowing for seamless integration with higher-precision data types used in deep learning frameworks.

Tensor Cores operate on 4x4 matrix tiles, performing matrix multiplication and accumulation (MMA) operations. These operations are highly parallel and can be efficiently executed on Tensor Cores. By exploiting the parallelism of Tensor Cores and their ability to process multiple elements simultaneously, significant performance gains can be achieved compared to traditional GPU implementations.

B. Matrix Multiplication Accumulate (MMA)

Matrix Multiplication Accumulate (MMA) is a key operation in many computational tasks, including deep learning, scientific simulations, and image processing. It involves three matrices: two input matrices (A and B) and an output matrix (C). The operation computes the matrix product of A and B and accumulates the result into C.

Mathematically, the MMA operation can be defined as follows:

$$C \leftarrow \alpha \cdot A \cdot B + \beta \cdot C$$

where A is an m x k matrix, B is a k x N matrix, C is an m x N matrix, and α and β are scalar coefficients.

The MMA operation is computationally intensive and can benefit greatly from hardware acceleration. Tensor Cores, with their ability to efficiently perform MMA operations, offer a promising solution for accelerating this operation.

III. EXPERIMENTAL SETUP

In this section, we describe the experimental setup used to evaluate Tensor Cores and compare their performance with CPU and GPU implementations of MMA.

A. Hardware Configuration

The experiments were conducted on a system equipped with an NVIDIA GPU that supports Tensor Cores. The specific GPU model is seen below.

TABLE I
GPU SPECIFICATIONS

GPU Model	NVIDIA Quadro 4000 RTX
Shared memory per threadblock	64 Kb
RAM	8 GB

B. Software Configuration

The experiments were conducted using the following software tools and frameworks:

- CUDA Toolkit version 11.8
- WMMA
- C++17 compiled with g++ 11.3.0

We implemented the MMA operation using different implementations: CPU, GPU (without Tensor Cores), and GPU (with Tensor Cores). The CPU implementation was written in C++, while the GPU implementations were developed using CUDA.

IV. CODE

In this section we present the main part of the code. The code shown is the computational part. Since this is deemed most important.

A. C++ code

This is the optimized C++ code
TS == tile size
N == matrix size
A, B and C are aligned matrices of size N*N for AVX 512, consting floatpoint32

```cpp
void MMA(float out, float A, float B, float
    C, int N){
float tile[TS][TS];
#pragma omp parallel private(tile){
 int ih, jh, kh, il, kl, jl;
#pragma omp simd aligned(A:32) aligned(B:32)
    aligned(out:32) collapse(3)
for (int ih = 0; ih < N; ih+=TS)
  for (int jh = 0; jh < N; jh+=TS)
   for (int kh = 0; kh < N; kh+=TS){
    for (int il = 0; il < TS; il ++)
     #pragma unroll
     for (int jl = 0; jl < TS; jl ++)
      tile[il][jl] = 0.;
     for (int il = 0; il < TS; il ++)
      for (int kl = 0; kl < TS; kl ++)
       #pragma unroll
       for (int jl = 0; jl < TS; jl ++)
        tile[il][jl] += A[(ih+il)*N+kh+kl]
            * B[(kh+kl)*N+jh+jl];
     for (int il = 0; il < TS; il ++)
      #pragma unroll
      for (int jl = 0; jl < TS; jl ++)
       out[(ih+il)*N+jh+jl] +=
           tile[il][jl];}}
#pragma omp parallel{
#pragma omp simd aligned(C:32)
    aligned(out:32) collapse(2)
 for ( int i = 0; i < N; i++)
  for ( int j = 0; j < N; j++)
   out[i*N+j] += C[i*N+j];}}
```

Text after it ...
Some notes about something and noce noice.

B. Cuda code

A, B and C are aligned floatingpoint 32 matrices of size N*N a_shared == locally shared part of the matrix
b_shared == locally shared part of the matrix
To make the code more compact, these simplifications are made:
Bs = Blocksize
tIdx = threadIdx
bIdx = blockIdxx
bDim = blockDimx

```
__global__ void MMA(float *out, float *A,
    float *B, float *C, int N){
__shared__ float a_shared [Bs][Bs];
__shared__ float b_shared [Bs][Bs];
// The kernel/thread global id
int row = bIdx.y * bDim.y + tIdx.y;
int col = bIdx.x * bDim.x + tIdx.x;
// Loop over the tiles
for (int tileNum = 0; tileNum < N/Bs;
    tileNum++){
  // j is the column index of the left matrix
  int j = tileNum*Bs + tIdx.x;
  int i = tileNum*Bs + tIdx.y;
  // load into shared memory, coalesced
  a_shared[tIdx.y][tIdx.x] = A[row*N + j];
  b_shared[tIdx.y][tIdx.x] = B[i*N + col];
  // sync before computation
  __syncthreads();
  for (int k = 0; k < Bs; k++){
    C[row*N + col] += a_shared[tIdx.y][k] *
        b_shared[k][tIdx.x];}
    __syncthreads();
    out[row*N + col] = C[row*N + col];}}
```

text yada yada

C. Tensor code

To make the code more compact, these simplifications are made:
Bs = Blocksize
tIdx = threadIdx
bIdx = blockIdxx
bDim = blockDimx

```
__global__ void mat_mul_add_tensor(half *a,
    half *b, float *C, float *d, int N){
// Tile using a 2D grid
int WarpX = (bIdx.x * bDim.x + tIdx.x) /
    warpSize;
int WarpY = (bIdx.y * bDim.y + tIdx.y);
// Declare the fragments
wmma::fragment<wmma::matrix_a, WMMA_M,
    WMMA_N, WMMA_K, half, wmma::row_major>
    a_frag;
wmma::fragment<wmma::matrix_b, WMMA_M,
    WMMA_N, WMMA_K, half, wmma::row_major>
    b_frag;
wmma::fragment<wmma::accumulator, WMMA_M,
    WMMA_N, WMMA_K, float> c_frag;
if(C != nullptr)
wmma::load_matrix_sync(c_frag, C + WarpX * 16
    + WarpY * 16 * N, N, wmma::mem_row_major);
```

```
else
wmma::fill_fragment(c_frag, 0.0f);
// Loop over k
for (int i = 0; i < N; i += WMMA_K){
  wmma::load_matrix_sync(a_frag, a + WarpY *
      16 * N + i, N);
  wmma::load_matrix_sync(b_frag, b + WarpX *
      16 + i * N, N);
  // Perform the matrix multiplication
  wmma::mma_sync(c_frag, a_frag, b_frag,
      c_frag);}
int cRow = WarpX * WMMA_M;
int cCol = WarpY * WMMA_N;
wmma::store_matrix_sync(d + cCol + cRow * N,
    c_frag, N, wmma::mem_row_major);}
```

V. BENCHMARKING METHODOLOGY

To evaluate the performance of Tensor Cores, we performed a series of experiments using different matrix sizes and precision settings. For each experiment, we measured the execution time of the MMA operation for CPU, GPU (without Tensor Cores), and GPU (with Tensor Cores) implementations.

We varied the matrix sizes from small to large to analyze the impact of matrix dimensions on performance. Additionally, we tested different precision settings, including full precision (FP32) and mixed precision (FP16).

To ensure accurate measurements, we repeated each experiment multiple times and calculated the average execution time. We also recorded the peak memory usage during each experiment to analyze the memory requirements of Tensor Cores.

VI. RESULTS AND ANALYSIS

In this section, we present the results of our experiments and analyze the performance of Tensor Cores compared to CPU and GPU implementations of MMA.

A. Performance Comparison

Insert Picture here
shows the execution time of the MMA operation for different matrix sizes and precision settings.

From the results, we can observe that the GPU implementation without Tensor Cores outperforms the CPU implementation for all matrix sizes. This is expected, as GPUs are highly parallel processors and can efficiently perform matrix computations. However, the GPU implementation with Tensor Cores significantly outperforms both the CPU and GPU implementations without Tensor Cores, demonstrating the superior performance of Tensor Cores for MMA operations.

We also observe that the performance of Tensor Cores is influenced by the matrix size. As the matrix size increases, the speedup achieved by Tensor Cores becomes more pronounced. This is because larger matrices can exploit the parallelism of Tensor Cores more effectively, resulting in higher throughput.

B. Impact of Precision

Next, we analyze the impact of precision on the performance of Tensor Cores. Figure

The results show that using mixed precision (FP16) with Tensor Cores can provide significant performance improvements compared to full precision (FP32). This is because Tensor Cores can process a larger number of elements simultaneously in FP16, leading to higher throughput. However, this comes to a cost as can be seen, as the matrix size grows, so does the error. Using FP16 then introduces a degree of numerical approximation, affecting the accuracy of the computation. Therefore, the choice of computing should be carefully considered based on the specific requirements of the application. Chosing cuda for high precision and tensorcores for faster times.

C. Optimization Techniques

To maximize the utilization of Tensor Cores and achieve optimal performance, we explored several optimization techniques, including:

- Matrix tiling: By partitioning the matrices into smaller tiles that fit the Tensor Core dimensions, we can improve data locality and exploit the parallelism of Tensor Cores more effectively.
- Memory access patterns: Optimizing memory access patterns can minimize data dependencies and ensure efficient memory utilization, reducing memory stalls and improving performance.
- Kernel fusion: Combining multiple operations into a single kernel can eliminate redundant memory accesses and kernel launches, reducing overhead and improving overall performance.

VII. LIMITATIONS AND CHALLENGES

The main limitations faced in this project was the limited information regarding how to program tensor cores. The cublass library, which is an optimized version of MMA with Tensor cores, developed by NVIDIA is not open source, so there is no real way of obtaining it. This led to experimenting with different setups to achieve the fastest copmutation possible. Another challenge to achieve fast computation for larger matrices with tensors are the fact that the GPU only has 64Kb of memory for each work group. With the tensors, the sum of shared memory is then N*N *(16+16). Which corresponds to the matrix size, and the byte size of the elements in the A and B matrix. This entails that the largest matrices that can fit this is of 32 x 32 elements of half precision.

VIII. CONCLUSION

In this paper, we have explored the performance of NVIDIA Tensor Cores for matrix multiplication and accumulation operations. We conducted experiments to compare the performance of Tensor Cores with CPU and GPU implementations, considering different matrix sizes and precision settings.

REFERENCES

[1] S. Markidis, S. W. D. Chien, E. Laure, I. B. Peng, and J. S. Vetter, "Nvidia tensor core programmability, performance precision," in *2018 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW)*, 2018, pp. 522–531.