

Evaluating tensor cores

John Carlsson, Cyprien Toulon de Courtex
Department of Computer Science
University of Salerno
Italy

Abstract— This project was aimed at evaluating tensor cores in cuda through experimnts to get a better understanding off high performance computing through different methods. With the ultimate goal to compare the time to calculate a matrix multiplication addition. At the start of the project, we were fairly certain that the fastest algorithm would be the tensor cores, followed by Cuda with shared memory,Cuda global memory, and at the bottom, an optimized CPU version. However, this was not the case. The results show that Cuda shared memory is the fastest with global memory just behind. For small matrices tensors appear superior.

I. INTRODUCTION

Over the past decade, the increasing demand for deep learning and AI applications has fueled the need for highly efficient tensor operations. Traditional GPU architectures are powerful but are often limited in their ability to fully exploit the potential of tensor computations due to the reliance on higher-precision floating-point arithmetic. Tensor cores address this limitation by leveraging mixed-precision arithmetic, combining high throughput with reduced precision.

Tensor cores, introduced in NVIDIA’s Volta architecture and further enhanced in subsequent architectures such as Turing and Ampere, have revolutionized the performance of tensor computations on GPUs. These specialized hardware units offer significant speedups by providing native support for low-precision tensor operations, specifically matrix multiplications and convolutions. The purpose of this project report is to present an evaluation of tensor cores, exploring their capabilities compared to CPU and GPU based computation.

In this project, we aim to evaluate tensor cores using the Matrix Multiplication Accumulate (MMA) with various workloads and compare the results to the CPU and GPU performances.

By conducting this comprehensive evaluation, we aim to provide insights into the capabilities and limitations of tensor cores. Additionally, the findings from this study will contribute to the broader understanding of GPU acceleration techniques for tensor operations and their potential impact on high-performance computing.

A. Motivation

The MMA operation plays a critical role in accelerating various computational tasks. It involves three matrices and is fundamental within many computational domains, such as deep learning, scientific simulations, and image processing. However, performing MMA efficiently and at scale is difficult due to the intensity of the operation.

The difficulties that arise when trying to implement a fast version of MMA are due to a combination of factors including memory access patterns, data dependencies, computational intensity, matrix sizes, hardware limitations, and optimization trade-offs.

The objective of this paper is to provide valuable insights for optimizing MMA computations and facilitate the adoption of tensor cores as a powerful tool for accelerating MMA operations in comparison to CPU and GPU. Our contributions will be as follows:

- We will measure the speedup achieved by tensor cores measured in actual seconds and explore how performance scales. with the size of the matrices.
- We will compare the performance of tensor cores with CPU and GPU implementations of MMA using different matrix sizes and precision settings.
- We will analyze the impact of data precision on the performance of tensor cores and investigate the trade-offs between accuracy and computation time.
- We will explore optimization techniques for maximizing the utilization of tensor cores and achieving optimal performance.
- We will provide insights into the limitations and challenges of tensor cores, including memory constraints, data dependencies, and scalability issues.

II. BACKGROUND

A. Tensor cores




Fig. 1. SM architecture.[1]

tensor cores are specialized hardware units introduced in NVIDIA GPUs that are designed to accelerate matrix com-

putations, specifically matrix multiplications and convolutions. They leverage mixed-precision arithmetic, combining high throughput with reduced precision to achieve significant speedups in tensor operations.

The key feature of tensor cores is their ability to perform mixed-precision operations using half-precision floating-point (FP16) data types. By using FP16, tensor cores can process a larger number of elements simultaneously, leading to improved throughput.

tensor cores operate on up to 16x16 matrix tiles, performing matrix multiplication and accumulation (MMA) operations. These operations are highly parallel and can be efficiently executed on tensor cores. By exploiting the parallelism of tensor cores and their ability to process multiple elements simultaneously, significant performance gains can be achieved compared to traditional GPU implementations.

The tensor operations are warp wise, this means that each tensor instructions must work on the same data within the same warp.

B. Matrix Multiplication Accumulate (MMA)

Matrix Multiplication Accumulate (MMA) is a key operation in many computational tasks, including deep learning, scientific simulations, and image processing. It involves three matrices: two input matrices (A and B) and an output matrix (C). The operation computes the matrix product of A and B and accumulates the result into C.

Mathematically, the MMA operation can be defined as follows:

$$C \leftarrow \alpha \cdot A \cdot B + \beta \cdot C$$

where A is an m x k matrix, B is a k x N matrix, C is an m x N matrix, and α and β are scalar coefficients.

The MMA operation is computationally intensive and can benefit greatly from hardware acceleration. tensor cores, with their ability to efficiently perform MMA operations, offer a promising solution for accelerating this operation.

III. EXPERIMENTAL SETUP

In this section, we describe the experimental setup used to evaluate tensor cores and compare their performance with CPU and GPU implementations of MMA. We did a multitude of tests with different parameters to see how performance scale over matrix size and how conducting many runs consequently affects performance over time.

We used square matrices with the sizes of 32, 64, 128, 256, 512, 1024, 2048, 4096 and 8192.

we conducted single comparison runs and then multiple runs to get a better average for time and error.

To understand the effect of each optimization methods, we produced different implementations, exploiting Cuda, shared memory, tensor cores, etc.

A. Hardware Configuration

The experiments were conducted on a system equipped with an NVIDIA GPU that supports tensor cores. The specific GPU model is seen below. The GPU belongs to a group of compute capability 7.5.

TABLE I
GPU SPECIFICATIONS[2]

GPU Model	NVIDIA Quadro 4000 RTX
Streamline Multiprocessors (SM)	4
Compute capability	7.5

B. Software Configuration

The experiments were conducted using the following software tools and frameworks:

- CUDA Toolkit version 11.8
- WMMA
- C++17 compiled with g++ 11.3.0

We implemented the MMA operation in different ways: CPU, GPU (without tensor cores), and GPU (with tensor cores). The CPU implementation was written in C++, while the GPU implementations were developed using CUDA.

IV. CODE

In this section we present the main part of the code. The code shown is the computational part. Since this is deemed most important. We show here the Cuda code and the tensor code.

A. CPU code

We implement a CPU version of the MMA operation to have a reference to verify our computation results. The CPU code is an implement of the MMA operation that uses the following optimizations:

- 1) tiling
- 2) parallelisation using OpenMP
- 3) compiler vectorisation for AVX512 using OpenMP
- 4) compiler optimization flags

B. Cuda code

The code is heavily simplified to keep the essence. For full code, see appendix.

```
__global__ void mat_mul_add(Out *out, In
    *A, In *B, Out *c, int n)
{
    sum_acc = c;
    for (tileNum = 0; tileNum <
        n/TILE_SIZE; tileNum++)
    { // load into shared memory, coalesced
        a_shared = A;
        b_shared = B;
        __syncthreads();
        for (int k = 0; k < TILE_SIZE; k++)
```

```

        {sum_acc += a_shared[threadIdx.y][k]
          * b_shared[k][threadIdx.x]; }
    }
    out = sum_acc;
}

```

C. Tensor code

To make the code more compact, these simplifications are made:

Bs = Blocksize

tIdx = threadIdx

bIdx = blockIdxx

bDim = blockDimx

```

__global__ void mat_mul_add_tensor(half *a,
    half *b, float *C, float *d, int N){
// Declare the fragments
wmma::fragment<wmma::matrix_a, WMMA_M,
    WMMA_N, WMMA_K, half, wmma::row_major>
    a_frag;
wmma::fragment<wmma::matrix_b, WMMA_M,
    WMMA_N, WMMA_K, half, wmma::row_major>
    b_frag;
wmma::fragment<wmma::accumulator, WMMA_M,
    WMMA_N, WMMA_K, float> c_frag;
wmma::load_matrix_sync(c_frag, C, N,
    wmma::mem_row_major);
// Loop over k
for (int i = 0; i < N; i += WMMA_K){
    wmma::load_matrix_sync(a_frag, a, N);
    wmma::load_matrix_sync(b_frag, b, N);
    // Perform the matrix multiplication
    wmma::mma_sync(c_frag, a_frag, b_frag,
        c_frag);}
    wmma::store_matrix_sync(d, c_frag, N,
        wmma::mem_row_major);}

```

V. BENCHMARKING METHODOLOGY

To evaluate the performance of tensor cores, we performed a series of experiments using different matrix sizes and precision settings. For each experiment, we measured the execution time of the MMA operation for CPU, GPU (without tensor cores), and GPU (with tensor cores) implementations.

We varied the matrix sizes from small to large to analyze the impact of matrix dimensions on performance. Additionally, we tested different precision settings, including full precision (FP32) and mixed precision (FP16).

To ensure accurate measurements, we repeated each experiment multiple times and calculated the average execution time. We also recorded the peak memory usage during each experiment to analyze the memory requirements of tensor cores.

VI. RESULTS AND ANALYSIS

In this section, we present the results of our experiments and analyze the performance of tensor cores compared to CPU and GPU implementations of MMA.

A. Performance Comparison

Insert Picture here
shows the execution time of the MMA operation for different matrix sizes and precision settings.




Fig. 2. Performance Comparison




Fig. 3. Performance Comparison total time with deallocation

ANALYSIS

B. Impact of Precision

Next, we analyze the impact of precision on the performance of tensor cores. Figure

The results show that using mixed precision (FP16) with tensor cores can provide significant performance improvements compared to full precision (FP32). This is because tensor cores can process a larger number of elements simultaneously in FP16, leading to higher throughput. This accuracy also improves with time for the matrix sizes. Using FP16

VII. LIMITATIONS AND CHALLENGES

The main limitations faced in this project was the limited information regarding how to program tensor cores. The cublass library, which is an optimized version of MMA with tensor cores, developed by NVIDIA is not open source, so there is no real way of obtaining it. This led to experimenting with different setups to achieve the fastest computation possible. Another challenge to achieve fast computation for larger matrices with tensors are the fact that the GPU only has 64Kb of memory for each work group. With the tensors, the sum of shared memory is then $N*N*(16+16)$. Which corresponds to the matrix size, and the byte size of the elements in the A and B matrix. This entails that the largest matrices that can fit this is of 32 x 32 elements of half precision.

VIII. CONCLUSION

In this paper, we have explored the performance of NVIDIA tensor cores for matrix multiplication and accumulation operations. We conducted experiments to compare the performance of tensor cores with CPU and GPU implementations, considering different matrix sizes and precision settings.

From our result we can see that tensor cores without any optimization have significantly better computational power than the other, optimized approaches. It is then our belief that to optimize the tensor core version would decrease the computational time further. The total time comparison shows that allocation of memory has a significant impact on the time from input to output. Therefore using tensor cores is optimal in problem where the same data is handled continuously, for example in convolutional neural network and image processing. The error seems to decrease with matrix size which can be attributed to what values we are using when testing. Using only integers, the error is 0.

REFERENCES

- [1] S. Markidis, S. W. D. Chien, E. Laure, I. B. Peng, and J. S. Vetter, "Nvidia tensor core programmability, performance & precision," in *2018 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW)*, 2018, pp. 522–531.
- [2] Nvidia, "Volta Tuning guide," <https://docs.nvidia.com/cuda/volta-tuning-guide/index.html>, 2017, (Accessed: 31 May 2023).

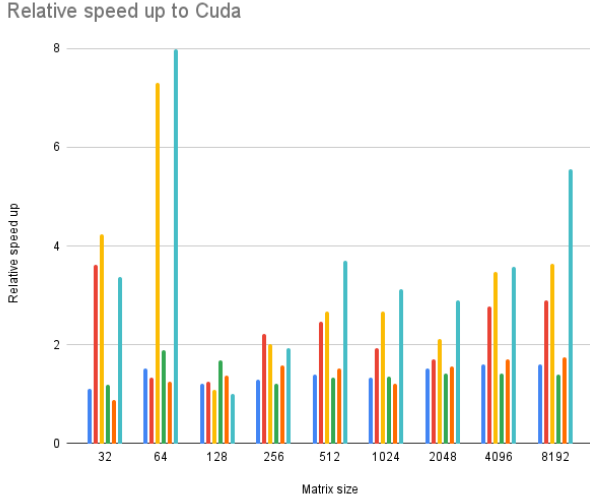


Fig. 4. Relative speedup to CUDA with global memory

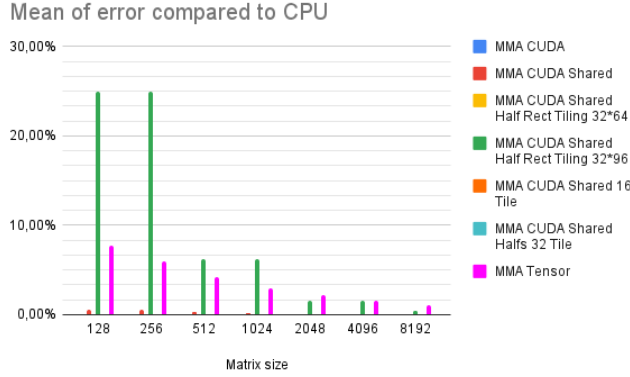


Fig. 5. Impact of Precision on tensor cores Performance

then introduces a degree of numerical approximation, affecting the accuracy of the computation. Therefore, the choice of computing should be carefully considered based on the specific requirements of the application.

C. Optimization Techniques

To maximize the utilization of our approaches and achieve optimal performance, we explored several optimization techniques, including:

- **Matrix tiling:** By partitioning the matrices into smaller tiles that fit the tensor Core dimensions, we can improve data locality and exploit the parallelism of tensor cores more effectively.
- **Memory access patterns:** Optimizing memory access patterns can minimize data dependencies and ensure efficient memory utilization, reducing memory stalls and improving performance.